

# Generative models

MIT

2022

# Generative and discriminative models

**Discriminative models**


Model:  $p(y|x)$ .

**Generative models**

Model:  $p(y, x)$ .

## Generative models:



- Generate datasets (when generation is a goal)
- Synthetic dataset generation (for train or fine-tuning)
- Latent dataset properties obtaining



# Data generation: example

Paganini et al., 2017:

- Model particle energy
- The modeling uses GAN
- Discrimination is done using GEANT software
- Result: good performance, generation is done 100-1000 times faster



# Data generation: example

Adams et al., 2010:

- The problem is to generate deep belief networks
- The model structure  $\Gamma$  is a sequence of adjacency matrices for each layer
- The generation is done using MCMC with Indian buffet process  $(\alpha, \beta)$  as a prior
- $\alpha, \beta$  can be interpreted as a width and sparsity of the structure




(a)  $\alpha = 1, \beta = 1$



(b)  $\alpha = \frac{1}{2}, \beta = 1$



(c)  $\alpha = 1, \beta = 2$



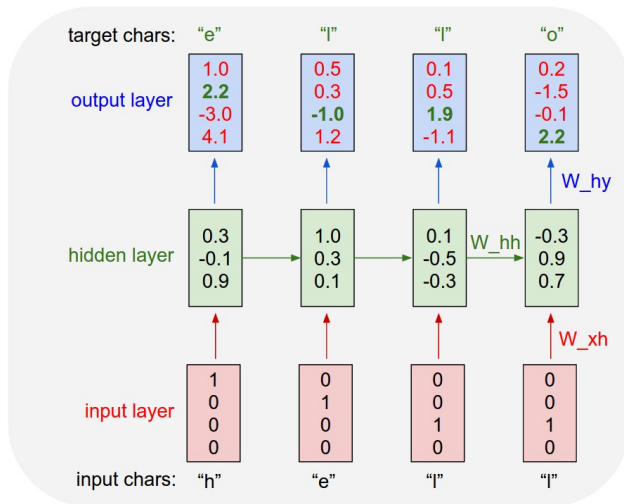
(d)  $\alpha = \frac{3}{2}, \beta = 1$

# How to build generative models?

- **Approach 1:** assign a likelihood function (“Fully-observed likelihood”), which decomposes object likelihood into parts (“Autoregressive models”).

# Example: CharRNN

Karpathy, 2015



# How to build generative models?

- **Approach 1:** assign a likelihood function (“Fully-observed likelihood”), which decomposes object likelihood into parts (“Autoregressive models”).

## **Problems:**

- ▶ hard to assign a proper likelihood function.
- ▶ computationally intensive inference.

# How to build generative models?

- **Approach 1:** assign a likelihood function (“Fully-observed likelihood”), which decomposes object likelihood into parts (“Autoregressive models”).

## **Problems:**

- ▶ hard to assign a proper likelihood function.
- ▶ computationally intensive inference.
- **Approach 2:** make an assumption that objects are generated by a latent variable, which is easier to analyze (“Latent variable models”).




## Example: autoencoder

Autoencoder is a model of dimension reduction:

$$H = \sigma(W_e X),$$

$$\|\sigma(W_d H) - X\|_2^2 \rightarrow \min.$$



# Autoencoder: generative model?

(Alain, Bengio 2012): consider regularized autoencoder:


$$\|f(x, \sigma) - x\|^2,$$

where  $\sigma$  is a noise level.

Then

$$\frac{\partial \log p(x)}{\partial x} = \frac{\|f(x, \sigma) - x\|^2}{\sigma^2} + o(1) \text{ при } \sigma \rightarrow 0.$$

Vector field induced by reconstruction error



# Variational autoencoder

Let the objects  $X$  be generated by latent variable  $h \sim \mathcal{N}(0, I)$ :

$$x \sim p(x|h, w).$$

$p(h|x, w)$  is unknown.

Maximize ELBO:


$$\log p(x|w) \geq \mathbb{E}_{q_\phi(h|x)} \log p(x|h, w) - D_{\text{KL}}(q_\phi(h|x) || p(h)) \rightarrow \max.$$

Distributions  $q_\phi(h|x)$  и  $p(x|h, w)$  are modeled by neural networks:

$$q_\phi(h|x) \sim \mathcal{N}(\mu_\phi(x), \sigma_\phi^2(x)),$$

$$p(x|h, w) \sim \mathcal{N}(\mu_w(h), \sigma_w^2(h)),$$


where  $\mu, \sigma$  are neural network's outputs.



# Variational autoencoder: generation process



(a) Learned Frey Face manifold



(b) Learned MNIST manifold

# How to build generative models?

- **Approach 1:** assign a likelihood function (“Fully-observed likelihood”), which decomposes object likelihood into parts (“Autoregressive models”).

## Problems:

- ▶ hard to assign a proper likelihood function.
- ▶ computationally intensive inference.
- **Approach 2:** make an assumption that objects are generated by a latent variable, which is easier to analyze (“Latent variable models”).

## Problems:

- ▶  $p(x)$  is intractible
- Problem of both methods: high likelihood and high sampling quality can be not independent (Theis et al., 2015).
- Given a noisy mixture:

$$p_w(x) = 0.01p_{\text{data}}(x) + 0.99p_{\text{noise}}(x), \log p_w(x) \geq \log p_{\text{data}}(x) - \log 100$$

- For another direction: overfitting

# How to build generative models?

- **Approach 1:** assign a likelihood function (“Fully-observed likelihood”), which decomposes object likelihood into parts (“Autoregressive models”).

## **Problems:**

- ▶ hard to assign a proper likelihood function.
  - ▶ computationally intensive inference.
- **Approach 2:** make an assumption that objects are generated by a latent variable, which is easier to analyze (“Latent variable models”).
- **Approach 3:** do not use likelihood and work straightforwardly with generative process (from likelihood modeling to statistical testing).

# Generative-adversarial models (Goodfellow et al., 2014)

**Main idea:** train two models, generator  $G$  and discriminator  $D$ :

$$\min_{W_G} \max_{w_D} \mathbb{E}_{x \in \mathcal{D}} \log p(x|w_D, D) + \mathbb{E}_{x \in p_G} \log(1 - p(x|w_D, D)).$$

The algorithm is iterative


- $\mathbb{E}_{x \in \mathcal{D}} \log p(x|w_D, D) \rightarrow \max_{w_D}$
- $\mathbb{E}_{x \in p_G} \log(1 - p(x|w_D, D)) \rightarrow \min_{w_G}$
- Alternative:  $\mathbb{E}_{x \in p_G} \log p(x|w_D, D) \rightarrow \max_{w_G}$

# GAN: optimality

When a discriminator is in global optimum, the generator minimizes  $JS$ :

$$-\log(4) + KL\left(p(x) \middle| \frac{p(x) + p_G(x)}{2}\right) + KL\left(p_G(x) \middle| \frac{p(x) + p_G(x)}{2}\right) \rightarrow \min_{w_G}.$$


**Consequent:** the optimal generator distribution:  $p_G = p(x)$ .





# Optimization details for GAN

- Generator optimization can be made in two regimes:  $E_{x \in p_G} \log(1 - p(x|w_D, D)) \rightarrow \min_{w_G}$  or  $E_{x \in p_G} \log p(x|w_D, D) \rightarrow \max_{w_G}$ : the optima coincide, but for the first regime the gradient is more smooth.
- Generator can converge to a local optimum and generate only similar objects (mode collapse).



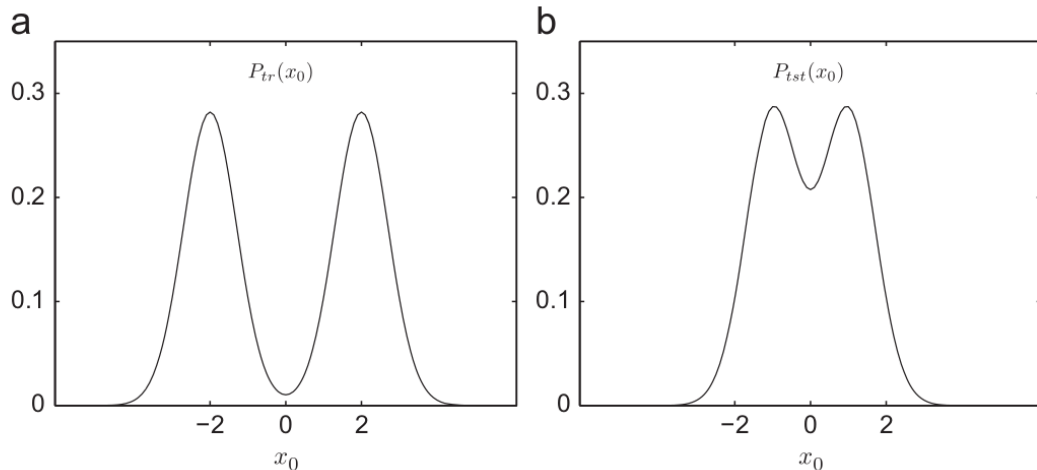
<https://machinelearningmastery.com/practical-guide-to-gan-failure-modes/>

# Dataset shift

Dataset shift is an event when distribution  $p(X, y)$  significantly differ for the training and test/inference phases.


- Covariate shift — difference in  $p(X)$
- Prior probability shift — difference in  $p(y)$
- Concept shift — difference in  $p(y|X)$

# Dataset shift




**Fig. 1.** Covariate shift:  $P_{tst}(y|x_0) = P_{tr}(y|x_0)$  and  $P_{tr}(x_0) \neq P_{tst}(x_0)$ . (a) Training data and (b) test data.

# Dataset shift



# Dataset shift



Moreno-Torres et al., 2012

# Evidence vs Cross-validation

Evidence:

$$\log p(X|f) = \log p(x_1|f) + \log p(x_2|x_1, f) + \cdots + \log p(x_n|x_1, \dots, x_{n-1}, f).$$



Leave-one-out:

$$\text{LOU} = E \log p(x_n|x_1, \dots, x_{n-1}, f).$$

Cross-validation uses mean value of the last term  $p(x_n|x_1, \dots, x_{n-1}, f)$  for complexity estimation.

Evidence considers **full** complexity.

# Evidence vs Cross-validation: example



# References

- Bishop C. M. Pattern recognition //Machine learning. – 2006. – T. 128. – №. 9.
- Paganini M., de Oliveira L., Nachman B. Accelerating science with generative adversarial networks: an application to 3D particle showers in multilayer calorimeters //Physical review letters. – 2018. – T. 120. – №. 4. – C. 042003.
- Antoran J., Miguel A. Disentangling and learning robust representations with natural clustering //2019 18th IEEE International Conference On Machine Learning And Applications (ICMLA). – IEEE, 2019. – C. 694-699.
- Adams R. P., Wallach H., Ghahramani Z. Learning the structure of deep sparse graphical models //Proceedings of the thirteenth international conference on artificial intelligence and statistics. – JMLR Workshop and Conference Proceedings, 2010. – C. 1-8.
- Alain G., Bengio Y. What regularized auto-encoders learn from the data-generating distribution //The Journal of Machine Learning Research. – 2014. – T. 15. – №. 1. – C. 3563-3593.
- Theis L., Oord A., Bethge M. A note on the evaluation of generative models //arXiv preprint arXiv:1511.01844. – 2015.
- Kingma D. P., Welling M. Auto-Encoding Variational Bayes //stat. – 2014. – T. 1050. – C. 10.
- Efron B., Tibshirani R. *An Introduction to the Bootstrap*, 1993.
- Moreno-Torres J. G. et al. A unifying view on dataset shift in classification //Pattern recognition. – 2012. – T. 45. – №. 1. – C. 521-530.
- Bakhteev O. Y., Strijov V. V. Comprehensive analysis of gradient-based hyperparameter optimization algorithms //Annals of Operations Research. – 2020. – T. 289. – №. 1. – C. 51-65.



# References

- Aditya Grover et al., Deep Generative Models tutorial, 2018: [goo.gl/H1prjP](https://goo.gl/H1prjP)
- Fei-Fei Li et al., Generative Models tutorial, 2017, [http://cs231n.stanford.edu/slides/2017/cs231n\\_2017\\_lecture13.pdf](http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture13.pdf)
- Shakir Mohamed et al., UAI 2017 Tutorial, 2017, <https://www.youtube.com/watch?v=JrO5fSsklSY>
- Andrej Karpathy: The Unreasonable Effectiveness of Recurrent Neural Networks, 2015: <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>
- Maxim Panov: Uncertainty, Out-of-distribution detection for NNs: [https://www.youtube.com/watch?v=N-p\\_qSLzoAI](https://www.youtube.com/watch?v=N-p_qSLzoAI)
- <https://machinelearningmastery.com/practical-guide-to-gan-failure-modes/>
- Cat generator: <https://github.com/aleju/cat-generator>